



The control layer for AI agents that take **real-world action**.

Before an agent moves money, changes access, or touches production, Sentinel judges its risk in real time. Routine actions pass. Risky, irreversible, or out-of-scope ones are **frozen and escalated to a human** — with a tamper-evident audit trail, in seconds.

• Live demo

`sentinel-sable-nu.vercel.app`

01 · CHALLENGE

Trust & permissions for autonomous agents

THEME A · STATEMENT A-1

"How might we design **permission layers, audit trails, and fail-safes** so users can trust agents to act on their behalf?"

AI agents can finally take actions. The question is no longer *can they* — it's **can we trust them to.**

02 · TARGET USER & USE CASE

Risk, fraud & compliance leaders at fintechs & digital banks

THE MOMENT

...at the exact moment an AI agent is about to **execute a consequential action** — a \$25,000 wire, a refund outside policy, a customer-data export — and they have no safe way to let it act.

INDUSTRY

Financial services

+ ROLE

Risk & compliance

+ MOMENT

An agent about to act

+ STRUGGLE

Can't safely delegate it

03 · PAIN & BOTTLENECK

PAIN — THE SYMPTOM

Capable agents sit unused. Teams won't let an agent move money, because one wrong, **irreversible** action — and a regulator asks why.

BOTTLENECK — THE ROOT CAUSE

Agents inherit broad, standing credentials with **no per-action permission model** — no scoped least-privilege, no human checkpoint, no audit. There's no safe way to delegate a single action.

Our solution attacks the **bottleneck** — a permission model for *actions*, not just identities.

04 · CURRENT ALTERNATIVE

What teams do today — and why it fails

OPTION A

Keep a human on everything

Safe, but it defeats the entire point of deploying an agent. Doesn't scale.

OPTION B

"Add a flag in a sprint"

Brittle and home-grown. Doesn't survive an audit — or the liability.

OPTION C

Trust each vendor's guardrails

Not neutral, not auditable. Each vendor grades its own homework.

An audit-grade control layer between the agent and the action

We help risk & compliance teams who struggle with trusting AI agents to act on money, at the moment an agent executes a consequential action, by solving the missing per-action permission model through a layer that scopes, judges, and escalates every action — so they can finally deploy their agents safely.

Deploy agents safely

Routine actions pass straight through, automatically.

Catch the risky action


Frozen and escalated to a human before it executes.

Prove it to a regulator

Every decision logged, tamper-evident, exportable.

06 · HOW IT WORKS

A hybrid engine — not a wrapper



Rules catch the obvious. The AI catches the fraud they miss.

LIVE EXAMPLE

A \$9,800 payment — just under the \$10k approval line, established vendor, reversible. It passes *every* deterministic rule. **gpt-5.4** catches it anyway: threshold-shopping + recently-changed bank details + urgency = invoice-redirection fraud. ● Caught by the risk model

Three catch types

Scope · guardrail · risk model
— each labelled.

Delegation & revocation

User→agent→action; cut off an
agent instantly.

Tamper-evident audit


SHA-256 hash-chained,
exportable log.

Visible & provable

Watch it reason live; edit a
rule, verdict changes.

08 · WHY NOW & DEFENSIBILITY

Finance-first. Horizontal by design.



The same engine already governs [DevOps](#), [identity](#), [healthcare](#), [support](#) — a deleted database or an admin-grant to an external account is caught exactly like a poisoned wire.



Build vs buy: a flag is a sprint. An audit-grade, cross-vendor control layer a regulator will accept is not — and it isn't the agent vendors' core product.

09 · LIVE TODAY

Everyone's building the agent. We build the layer that makes it **deployable**.

Working engine

Two-layer hybrid + live gpt-5.4 verdicts, deployed and running.

Provable, not claimed

Watch it reason; edit a rule and the verdict changes, live.

Audit-ready

Every decision hash-chained and exportable.

• Try it live

`sentinel-sable-nu.vercel.app`